1 Claim

100 AVeriTeC claims (50 Supported, 50 Refuted)
claim, date and speaker

2 LLM

closed-book, temperature 0: Gemini, Claude, GPT
reasoning + label · 300 outputs, 32 wrong

3 Decompose

GPT-6.1 splits the reasoning into facts
1,612 facts + 100 claims to check

4 Evidence

per fact and per claim

Gold
annotators' QA pairs

BM25
top 5, word match

Similarity
top 5, embeddings

Verifier-guided
NLI picks top 5 of 20

retrieval pool: 10,498 web sentences

5 Verify

DeBERTa-v3 NLI, same in every condition
entailment / neutral / contradiction

6 Flag

flag if any fact is refuted, or if the claim
verdict opposes the LLM label

7 Evaluate

flags vs. wrong outputs
AUROC with 95% bootstrap CI over claims

gold label: used only in step 7

LLM label (saved in step 2)